

Delta Lake and Delta Tables

Delta Lake and Delta Tables

Introduction to Delta Lake and Delta Tables

Working with Delta Tables

Versioning and Time Travel

Schema Validation

Delta Table Internals

Optimizations



Databricks Cloud



Delta Lake

What is Delta Lake?

Delta Lake is an open-source storage framework.

Key Features

- ACID transactions
- Deletes, update, and merge
- Schema enforcement and evolution
- Data versioning and Time travel
- Streaming and batch unification
- Optimizations

How to use delta lake?

Delta Lake is available as a connector



Apache Spark™

Spark

This connector allows Apache Spark™ to read from and write to Delta Lake.



Snowflake (Beta)

This preview allows Snowflake to read from Delta Lake via an external table.



Redshift

Redshift AWS manifest

This utility allows AWS Redshift to read from Delta Lake using a manifest file.



Apache Hive

Hive standalone

This connector allows Apache Hive to read from Delta Lake.



PrestoDB

PrestoDB standalone

This connector allows PrestoDB to read from Delta Lake.



Apache Flink

Flink standalone

This connector allows Apache Flink to write to Delta Lake.



Power BI

PowerBI community

This connector allows Power BI to read from Delta Lake.



Kafka Delta Ingest

Kafka Rust

This project builds a highly efficient daemon for streaming data through Apache Kafka into Delta Lake.



Apache Beam (Beta)

Beam standalone community

This connector allows Apache Beam to read from Delta Lake.